

Model Development Phase

Date	14 July 2024
Team ID	SWTID1720085076
Project Title	Rice Type Classification using CNN
Maximum Marks	10 Marks

Initial Model Training Code, Model Validation and Evaluation Report

The initial model training code will be showcased in the future through a screenshot. The model validation and evaluation report will include a summary and training and validation performance metrics for multiple models, presented through respective screenshots.

Initial Model Training Code (5 marks):

```
✓ [14] mobile_net_url = 'https://tfhub.dev/google/tf2-preview/mobilenet_v2/feature_vector/4'
      mobile_net = hub.KerasLayer(mobile_net_url, input_shape=(224, 224, 3), trainable=True)

✓ # Build the model
  num_labels = 5
  model = keras.Sequential([
    data_augmentation,
    mobile_net,
    keras.layers.Dense(128, activation='relu'),
    keras.layers.BatchNormalization(),
    keras.layers.Dropout(0.5),
    keras.layers.Dense(num_labels, activation='softmax')
  ])
```

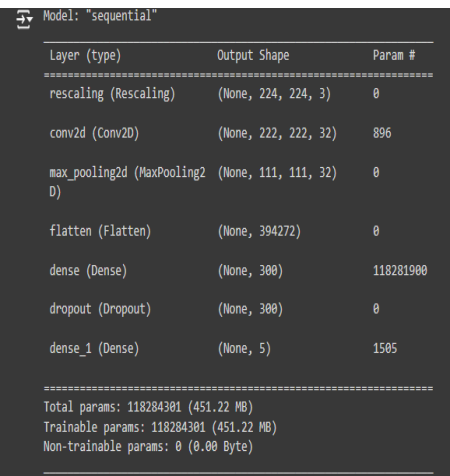
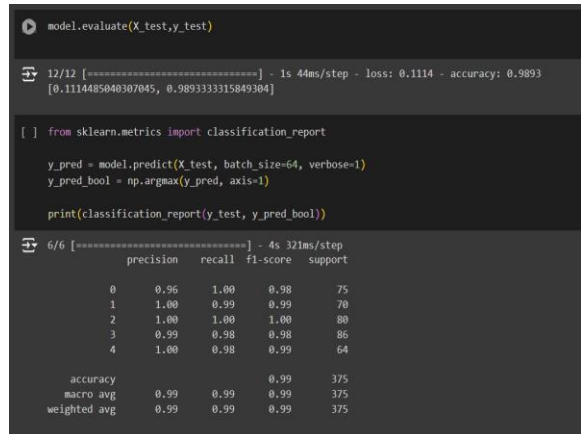
```
✓ [16] # Compile the model
      model.compile(
          optimizer=keras.optimizers.Adam(learning_rate=1e-4),
          loss=tf.keras.losses.SparseCategoricalCrossentropy(),
          metrics=['accuracy']
      )
```

```
✓ [17] early_stopping = keras.callbacks.EarlyStopping(monitor='val_loss', patience=5, restore_best_weights=True)
      reduce_lr = keras.callbacks.ReduceLROnPlateau(monitor='val_loss', factor=0.2, patience=3, min_lr=1e-6)
```

```
2m ▶ history = model.fit(X_train, y_train, epochs=10, validation_data=(X_val, y_val), callbacks=[early_stopping, reduce_lr])

Epoch 1/10
55/55 [=====] - 50s 261ms/step - loss: 0.6400 - accuracy: 0.8114 - val_loss: 1.1510 - val_accuracy: 0.5653 - lr: 1.0000e-04
Epoch 2/10
55/55 [=====] - 8s 148ms/step - loss: 0.1877 - accuracy: 0.9680 - val_loss: 0.9858 - val_accuracy: 0.6107 - lr: 1.0000e-04
Epoch 3/10
55/55 [=====] - 8s 151ms/step - loss: 0.1424 - accuracy: 0.9823 - val_loss: 0.4330 - val_accuracy: 0.8880 - lr: 1.0000e-04
Epoch 4/10
55/55 [=====] - 8s 151ms/step - loss: 0.1379 - accuracy: 0.9817 - val_loss: 0.2333 - val_accuracy: 0.9627 - lr: 1.0000e-04
Epoch 5/10
55/55 [=====] - 8s 149ms/step - loss: 0.1265 - accuracy: 0.9851 - val_loss: 0.2738 - val_accuracy: 0.9200 - lr: 1.0000e-04
Epoch 6/10
55/55 [=====] - 9s 155ms/step - loss: 0.1113 - accuracy: 0.9926 - val_loss: 0.2110 - val_accuracy: 0.9440 - lr: 1.0000e-04
Epoch 7/10
55/55 [=====] - 8s 151ms/step - loss: 0.1053 - accuracy: 0.9926 - val_loss: 0.1534 - val_accuracy: 0.9680 - lr: 1.0000e-04
Epoch 8/10
55/55 [=====] - 8s 147ms/step - loss: 0.1218 - accuracy: 0.9874 - val_loss: 0.1722 - val_accuracy: 0.9627 - lr: 1.0000e-04
Epoch 9/10
55/55 [=====] - 8s 150ms/step - loss: 0.1182 - accuracy: 0.9886 - val_loss: 0.1024 - val_accuracy: 0.9947 - lr: 1.0000e-04
Epoch 10/10
55/55 [=====] - 8s 151ms/step - loss: 0.1091 - accuracy: 0.9920 - val_loss: 0.0920 - val_accuracy: 0.9973 - lr: 1.0000e-04
```

Model Validation and Evaluation Report (5 marks):

Model	Summary	Training and Validation Performance Metrics
Model 1 MobileNet	 <pre>Model: "sequential" Layer (type) Output Shape Param # ----- rescaling (Rescaling) (None, 224, 224, 3) 0 conv2d (Conv2D) (None, 222, 222, 32) 896 max_pooling2d (MaxPooling2D) (None, 111, 111, 32) 0 flatten (Flatten) (None, 394272) 0 dense (Dense) (None, 300) 118281900 dropout (Dropout) (None, 300) 0 dense_1 (Dense) (None, 5) 1505 ----- Total params: 118284301 (451.22 MB) Trainable params: 118284301 (451.22 MB) Non-trainable params: 0 (0.00 Byte)</pre>	 <pre>model.evaluate(X_test, y_test) 12/12 [=====] - 1s 44ms/step - loss: 0.1114 - accuracy: 0.9893 [0.1114485040307045, 0.9893333315849304] [] from sklearn.metrics import classification_report y_pred = model.predict(X_test, batch_size=64, verbose=1) y_pred_bool = np.argmax(y_pred, axis=1) print(classification_report(y_test, y_pred_bool)) 6/6 [=====] - 4s 321ms/step precision recall f1-score support 0 0.96 1.00 0.98 75 1 1.00 0.99 0.99 70 2 1.00 1.00 1.00 80 3 0.99 0.98 0.98 86 4 1.00 0.98 0.99 64 accuracy 0.99 375 macro avg 0.99 0.99 0.99 375 weighted avg 0.99 0.99 0.99 375</pre>

Model 2 ALEXNET

Model: "model"

Layer (type)	Output Shape	Param #
input_1 (InputLayer)	[(None, 240, 240, 3)]	0
conv2d (Conv2D)	(None, 58, 58, 96)	34944
batch_normalization (Batch Normalization)	(None, 58, 58, 96)	384
max_pooling2d (MaxPooling2D)	(None, 28, 28, 96)	0
conv2d_1 (Conv2D)	(None, 24, 24, 256)	614656
batch_normalization_1 (Batch Normalization)	(None, 24, 24, 256)	1024
max_pooling2d_1 (MaxPooling2D)	(None, 11, 11, 256)	0
conv2d_2 (Conv2D)	(None, 9, 9, 384)	885120
conv2d_3 (Conv2D)	(None, 7, 7, 384)	1327488
conv2d_4 (Conv2D)	(None, 5, 5, 256)	884992
max_pooling2d_2 (MaxPooling2D)	(None, 2, 2, 256)	0
flatten (Flatten)	(None, 1024)	0
dense (Dense)	(None, 4096)	4198400
dropout (Dropout)	(None, 4096)	0
dense_1 (Dense)	(None, 4096)	16781312
dropout_1 (Dropout)	(None, 4096)	0
dense_2 (Dense)	(None, 5)	20485

```
[ ] # Assuming 'Alex_model' is the history object from your model training
training_acc_alex = Alex_model.history['accuracy']
val_acc_alex = Alex_model.history['val_accuracy']

best_training_acc = max(training_acc_alex)
best_val_acc = max(val_acc_alex)

# Converting to percentage and rounding off to 2 decimal places
best_training_acc_percentage = round(best_training_acc * 100, 2)
best_val_acc_percentage = round(best_val_acc * 100, 2)

print("Best Training Accuracy: ", best_training_acc_percentage, "%")
print("Best Validation Accuracy: ", best_val_acc_percentage, "%")
```

Best Training Accuracy: 93.24 %
Best Validation Accuracy: 86.86 %

Model 3 VGGNET

Model: "model_1"

Layer (type)	Output Shape	Param #
input_2 (InputLayer)	[(None, 240, 240, 3)]	0
conv2d_5 (Conv2D)	(None, 238, 238, 64)	1792
conv2d_6 (Conv2D)	(None, 236, 236, 64)	36928
batch_normalization_2 (Batch Normalization)	(None, 236, 236, 64)	256
max_pooling2d_3 (MaxPooling2D)	(None, 118, 118, 64)	0
conv2d_7 (Conv2D)	(None, 116, 116, 128)	73856
conv2d_8 (Conv2D)	(None, 114, 114, 128)	147584
batch_normalization_3 (Batch Normalization)	(None, 114, 114, 128)	512
max_pooling2d_4 (MaxPooling2D)	(None, 57, 57, 128)	0
conv2d_9 (Conv2D)	(None, 55, 55, 256)	295168
conv2d_10 (Conv2D)	(None, 53, 53, 256)	590080

```
# Assuming 'VGG_model' is the history object from your model training
training_acc_alex = VGG_model.history['accuracy']
val_acc_alex = VGG_model.history['val_accuracy']

best_training_acc = max(training_acc_alex)
best_val_acc = max(val_acc_alex)

# Converting to percentage and rounding off to 2 decimal places
best_training_acc_percentage = round(best_training_acc * 100, 2)
best_val_acc_percentage = round(best_val_acc * 100, 2)

print("Best Training Accuracy: ", best_training_acc_percentage, "%")
print("Best Validation Accuracy: ", best_val_acc_percentage, "%")
```

Best Training Accuracy: 83.63 %
Best Validation Accuracy: 83.86 %